

Bonus Question (20 points)

Create a data quality report for the Auto-MPG dataset.

Provide the data quality tables, distributions of categorical and nominal variables.

Also provide your solutions for handling outliers and missing values.

Create the data quality tables after handling outliers and missing values.

Provide this as a separate PDF file. You can use the cells below to find statistics and create visualizations.

In [1]:

```
%matplotlib inline
import pandas as pd
import numpy as np
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
adf = pd.read_csv('auto-mpg.csv')
adf.head()
```

Out[2]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	carname
0	18.0	8	307.0	130.0	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150.0	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150.0	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140.0	3449	10.5	70	1	ford torino

Data Quality Report before handling missing values and outliers

In [3]:

```
class DQR():
    def __init__(self):

        self.dqr_cont = pd.DataFrame(columns=['Feature', 'Count', '% Miss', 'Cardinality', 'Min', '1st Q', 'Mean', 'Median', '3rd Q', 'Max', 'S.D.'])
        self.dqr_cat = pd.DataFrame(columns=['Feature', 'Count', '% Miss', 'Cardinality', 'Mode', 'Mode freq', 'Mode %', '2nd Mode', '2nd Mode freq', '2nd Mode %'])
        self.dqr_cont['Feature'] = ['mpg', 'displacement', 'horsepower', 'weight', 'acceleration', 'cylinders']
        self.dqr_cat['Feature'] = ['origin', 'carname', 'year']

    def report(self, data):
        data = pd.DataFrame(data)
        for col in data.columns:
            if (col == 'origin' or col == 'carname' or col == 'year'):
                mode = data[col].value_counts().index
                mode_freq = data[col].value_counts().values

                self.dqr_cat.loc[self.dqr_cat['Feature'] == col, 'Count'] = data[col].count()
                self.dqr_cat.loc[self.dqr_cat['Feature'] == col, 'Cardinality'] = data[col].nunique()

                self.dqr_cat.loc[self.dqr_cat['Feature'] == col, 'Mode'] = mode[0]
                self.dqr_cat.loc[self.dqr_cat['Feature'] == col, 'Mode freq'] = mode_freq[0]
                self.dqr_cat.loc[self.dqr_cat['Feature'] == col, 'Mode %'] = (mode_freq[0] / (data[col].count())) * 100
                self.dqr_cat.loc[self.dqr_cat['Feature'] == col, '2nd Mode'] = mode[1]
```

```

        self.dqr_cat.loc[self.dqr_cat['Feature'] == col, '2nd Mode'] = mode[1]
        self.dqr_cat.loc[self.dqr_cat['Feature'] == col, '2nd Mode freq'] = mode_freq[1]
        self.dqr_cat.loc[self.dqr_cat['Feature'] == col, '2nd Mode %'] = (mode_freq[1] /
(data[col].count()))*100
        self.dqr_cat.loc[self.dqr_cat['Feature'] == col, '% Miss'] = ((data[col].isnull().su
m())/ (data[col].count()))*100

    else:
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, 'Count'] = data[col].count()
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, '% Miss'] = ((data[col].isnull()
sum())/ (data[col].count()))*100
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, 'Cardinality'] =
data[col].nunique()
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, 'Min'] = data[col].min()
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, '1st Q'] =
data[col].quantile(0.25)
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, 'Mean'] = data[col].mean()
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, 'Median'] = data[col].median()
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, '3rd Q'] =
data[col].quantile(0.75)
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, 'Max'] = data[col].max()
        self.dqr_cont.loc[self.dqr_cont['Feature'] == col, 'S.D. ' = data[col].std()

obj1 = DQR()
obj1.report(adf.copy())
# print("Data Quality report for Categorical variable: \n")
# print(obj1.dqr_cat)
# print("Data Quality report for Continuous variable: \n")
# print(obj1.dqr_cont)

```

In [4]:

```

print("\n Data Quality report for Continuous variable:- \n")
obj1.dqr_cont.head(6)

```

Data Quality report for Continuous variable:-

Out[4]:

	Feature	Count	% Miss	Cardinality	Min	1st Q	Mean	Median	3rd Q	Max	S.D.
0	mpg	398	2.01005	129	9	17.5	23.5146	23	29	46.6	7.81598
1	displacement	406	0	83	68	105	194.78	151	302	455	104.922
2	horsepower	400	1.5	93	46	75.75	105.082	95	130	230	38.7688
3	weight	406	0	357	19	2220	2952.31	2811	3612	5140	891.587
4	acceleration	406	0	96	8	13.7	15.5197	15.5	17.175	24.8	2.80336
5	cylinders	406	0	6	3	4	5.5	4	8	16	1.78989

In [5]:

```

print("\n Data Quality report for Categorical variable:- \n")
obj1.dqr_cat.head()

```

Data Quality report for Categorical variable:-

Out[5]:

	Feature	Count	% Miss	Cardinality	Mode	Mode freq	Mode %	2nd Mode	2nd Mode freq	2nd Mode %
0	origin	406	0	3	1	254	62.5616	3	79	19.4581
1	carname	406	0	312	ford pinto	6	1.47783	toyota corolla	5	1.23153
2	year	406	0	13	73	40	9.85222	78	36	8.867

Handling outliers

In [6]:

```
# Answer to Q4 goes here
# print(adf.describe())

print(adf.describe())
print("Skew before :-----")
print(adf.skew())
def outlier_detection(col):
    median = col.median()
    std = col.std()
    min_range = median - 3*std
    max_range = median + 3*std
    outlier = (col[(col > max_range) | (col < min_range)])
    print("Outliers in ", col.name)
    print(outlier)
    col.loc[col > max_range] = (col[col < max_range]).max()
    col.loc[col < min_range] = (col[col > min_range]).min()
    return col

adf['mpg'] = outlier_detection(adf['mpg'].copy())
adf['displacement'] = outlier_detection(adf['displacement'].copy())
adf['horsepower'] = outlier_detection(adf['horsepower'].copy())
adf['acceleration'] = outlier_detection(adf['acceleration'].copy())
adf['weight'] = outlier_detection(adf['weight'].copy())
adf['cylinders'] = outlier_detection(adf['cylinders'].copy())

print("Skew after :-----")
print(adf.skew())
print(adf.describe())
```

	mpg	cylinders	displacement	horsepower	weight \
count	398.000000	406.000000	406.000000	400.000000	406.000000
mean	23.514573	5.500000	194.779557	105.082500	2952.305419
std	7.815984	1.789889	104.922458	38.768779	891.587329
min	9.000000	3.000000	68.000000	46.000000	19.000000
25%	17.500000	4.000000	105.000000	75.750000	2220.000000
50%	23.000000	4.000000	151.000000	95.000000	2811.000000
75%	29.000000	8.000000	302.000000	130.000000	3612.000000
max	46.600000	16.000000	455.000000	230.000000	5140.000000

	acceleration	year	origin
count	406.000000	406.000000	406.000000
mean	15.519704	75.921182	1.568966
std	2.803359	3.748737	0.797479
min	8.000000	70.000000	1.000000
25%	13.700000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.175000	79.000000	2.000000
max	24.800000	82.000000	3.000000

Skew before :-----

```
mpg          0.457066
cylinders     0.906124
displacement  0.694130
horsepower    1.034079
weight        0.163454
acceleration  0.230224
year          0.020912
origin        0.932399
dtype: float64
Outliers in mpg
329    46.6
Name: mpg, dtype: float64
Outliers in displacement
Series([], Name: displacement, dtype: float64)
Outliers in horsepower
6      220.0
7      215.0
8      225.0
19     225.0
```

```

31      215.0
101     215.0
102     225.0
123     230.0
Name: horsepower, dtype: float64
Outliers in acceleration
306     24.8
402     24.6
Name: acceleration, dtype: float64
Outliers in weight
194     42
227     19
344     22
398     26
Name: weight, dtype: int64
Outliers in cylinders
260     16
Name: cylinders, dtype: int64
Skew after :-----
mpg          0.445905
cylinders     0.501657
displacement  0.694130
horsepower    0.952266
weight        0.492936
acceleration  0.187489
year          0.020912
origin        0.932399
dtype: float64

```

	mpg	cylinders	displacement	horsepower	weight \
count	398.000000	406.000000	406.000000	400.000000	406.000000
mean	23.509548	5.480296	194.779557	104.857500	2967.928571
std	7.801735	1.716544	104.922458	38.111723	853.167769
min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.500000	4.000000	105.000000	75.750000	2220.000000
50%	23.000000	4.000000	151.000000	95.000000	2811.000000
75%	29.000000	8.000000	302.000000	130.000000	3612.000000
max	44.600000	8.000000	455.000000	210.000000	5140.000000

	acceleration	year	origin
count	406.000000	406.000000	406.000000
mean	15.514778	75.921182	1.568966
std	2.788013	3.748737	0.797479
min	8.000000	70.000000	1.000000
25%	13.700000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.175000	79.000000	2.000000
max	23.700000	82.000000	3.000000

Handling missing Values

In [7]:

```

# Answer to Q5 goes here
# your code ....
from sklearn.impute import KNNImputer
knn_imputer = KNNImputer(n_neighbors=3)
impute_copy = adf[['mpg', 'horsepower', 'cylinders', 'displacement', 'weight']].copy()
print((impute_copy.isnull()).sum())

adf_transformed = knn_imputer.fit_transform(impute_copy)

print(sum(np.isnan(adf_transformed)))

adf_trans = pd.DataFrame(index=range(adf.shape[0]), columns=['mpg', 'horsepower', 'cylinders', 'displacement', 'weight'])
adf_trans = pd.DataFrame(adf_transformed, dtype=None, copy=False, index = adf_trans.index, columns=adf_trans.columns)
adf_trans[['acceleration', 'year', 'origin', 'carname']] =
adf[['acceleration', 'year', 'origin', 'carname']]

```

```

mpg          8
horsepower    6
cylinders      0
displacement  0

```

```
weight          0
dtype: int64
[0 0 0 0 0]
```

In [8]:

```
adf_trans.head()
```

Out[8]:

	mpg	horsepower	cylinders	displacement	weight	acceleration	year	origin	carname
0	18.0	130.0	8.0	307.0	3504.0	12.0	70	1	chevrolet chevelle malibu
1	15.0	165.0	8.0	350.0	3693.0	11.5	70	1	buick skylark 320
2	18.0	150.0	8.0	318.0	3436.0	11.0	70	1	plymouth satellite
3	16.0	150.0	8.0	304.0	3433.0	12.0	70	1	amc rebel sst
4	17.0	140.0	8.0	302.0	3449.0	10.5	70	1	ford torino

Data Quality Report after handling outliers and missing values

In [9]:

```
obj2 = DQR()
obj2.report(adf_trans)
```

In [10]:

```
print("\n Data Quality report for Continuous variable:- \n")
obj2.dqr_cont.head(6)
```

Data Quality report for Continuous variable:-

Out[10]:

	Feature	Count	% Miss	Cardinality	Min	1st Q	Mean	Median	3rd Q	Max	S.D.
0	mpg	406	0	133	9	17	23.435	22.75	29	44.6	7.80953
1	displacement	406	0	83	68	105	194.78	151	302	455	104.922
2	horsepower	406	0	94	46	75	104.547	95	129	210	37.9754
3	weight	406	0	353	1613	2220	2967.93	2811	3612	5140	853.168
4	acceleration	406	0	94	8	13.7	15.5148	15.5	17.175	23.7	2.78801
5	cylinders	406	0	5	3	4	5.4803	4	8	8	1.71654

In [11]:

```
print("\n Data Quality report for Catgorical variable:----- \n")
obj2.dqr_cat.head()
```

Data Quality report for Catgorical variable:-----

Out[11]:

	Feature	Count	% Miss	Cardinality	Mode	Mode freq	Mode %	2nd Mode	2nd Mode freq	2nd Mode %
0	origin	406	0	3	1	254	62.5616	3	79	19.4581
1	carname	406	0	312	ford pinto	6	1.47783	toyota corolla	5	1.23153
2	carname	406	0	10	toyota corolla	10	2.46552	toyota corolla	20	4.92637

2	year	406	0	13	73	40	9.85222	78	36	8.867
Feature	Count	% Miss	Cardinality	Mode	Mode freq	Mode %	2nd Mode	2nd Mode	2nd Mode	2nd Mode

In [12]:

```
import seaborn as sns
sns.set()
sns.pairplot(adf_trans, hue='origin', diag_kind='hist', markers=['s', 'd', 'o'], height=2.0)
plt.show()
```

